

HerHealthEval: Do Language Models Understand Women’s Health Communication Across Languages and Cultures?

Anonymous ACL submission

Abstract

Large language models are increasingly used to answer women’s-health questions, but fluent answers do not establish that a model has correctly understood what a patient is trying to communicate. We present HerHealthEval, a preliminary multilingual evaluation framework and protocol for studying language understanding in women’s-health communication across languages and linguistic registers.

Our controlled evaluation suite covers menstruation, polycystic ovary syndrome (PCOS)/hormonal concerns, and fertility in English, Arabic, and French, pairing clinical, lay, indirect, ambiguous, and emotionally concerned expressions with predefined targets for concern category, risk level, recommended action, and clarification need. The protocol requires native-speaker validation of Arabic and French for meaning preservation, naturalness, and register before cross-lingual evaluation or adaptation, while evaluation cases remain separated from leakage-controlled adaptation data. The experimental design centers Qwen3.5-9B and specifies a general-purpose baseline, English domain adaptation, and subsequent Arabic/French adaptation, evaluated through misunderstanding, under-triage, clarification, unsafe-response, and cross-style/cross-language consistency measures. Rather than claiming comprehensive clinical coverage, this preliminary study establishes a failure-oriented protocol for testing women’s-health language understanding beyond answer generation.

1 Introduction

The central question here is not whether a language model can produce a plausible answer to a women’s-health question, but whether it first understands what the patient is trying to communicate. The same concern may be stated clinically (“I have irregular menstrual bleeding”), colloquially (“my period is all over the place”), or indirectly (“my monthly

problem has changed”). These expressions point to the same underlying concern but differ in explicitness, emotion, and cultural register; misreading that variation can lead a model to confuse the concern, understate urgency, offer false reassurance, or answer when it should first ask for missing information.

Most women’s-health LLM studies instead evaluate the generated answer rather than whether the input was understood, and multilingual health QA work shows the same question can elicit inconsistent information across languages. Together, these gaps motivate a controlled evaluation of *interpretation* rather than another answer-generation study; §2 reviews this literature in full.

We present **HerHealthEval**, an initial investigation of language understanding in multilingual women’s-health communication. The evaluation suite spans menstruation, PCOS/hormonal concerns, and fertility, holding each underlying case fixed while varying canonical, clinical, layperson, indirect/culturally sensitive, ambiguous, and emotionally concerned expression. Matched English cases are localized into Arabic and French under a native-speaker-approved protocol before entering cross-lingual evaluation or adaptation. Each item has predefined targets for concern category, risk level, recommended action, and clarification need, and the evaluation lineage is kept separate from the adaptation corpus by source- and paraphrase-family exclusion.

We organize the study around three research questions:

- RQ1:** How robust are LLM-based clinical reasoning systems to variation in linguistic register when interpreting semantically equivalent women’s health concerns?
- RQ2:** Do multilingual LLMs exhibit systematic differences in interpretation, triage calibra-

082	tion, and clarification behavior across English,	and Gély et al. (2026) compare LLM answers to	130
083	Arabic, and French?	infertility FAQs against specialist answers. Maurya	131
084	3. RQ3: What is the impact of domain-specific	et al. (2026) broaden coverage through WHBench,	132
085	and multilingual adaptation on linguistic ro-	an expert-crafted, safety-aware rubric. These stud-	133
086	bustness, uncertainty handling, and clarifica-	ies motivate domain-specific adaptation and careful	134
087	tion generation in women’s health symptom	response evaluation but treat the input question as	135
088	understanding?	fixed; none tests whether the same concern is in-	136
		terpreted consistently as its register or language	137
		changes.	138
089	The contributions align with these questions: (1)	2.2 Multilingual and Sociocultural	139
090	women’s-health interaction framed as a controlled	Understanding	140
091	language-understanding problem across register;	Jin et al. (2024) introduce X LingEval to assess cor-	141
092	(2) a multilingual evaluation protocol with matched	rectness, consistency, and verifiability on healthcare	142
093	language variants, a native-speaker validation gate,	questions across English, Spanish, Chinese, and	143
094	and leakage-separated adaptation data; and (3) a re-	Hindi, finding substantial cross-language dispari-	144
095	producible, failure-oriented comparison of general-	ties; subsequent work likewise reports inconsistent	145
096	purpose and sequentially adapted models on mis-	information across languages for equivalent ques-	146
097	understanding, triage, clarification, safety, and con-	tions (Schlicht et al., 2025). Foundational cross-	147
098	sistency. These contributions are deliberately pre-	cultural NLP work distinguishes linguistic diversity	148
099	liminary: an evaluation protocol and early failure	from cultural variation, identifying linguistic form,	149
100	patterns, not comprehensive coverage or clinical	common ground, topic relevance, and communica-	150
101	readiness.	tive objectives as separate dimensions (Hershcovich	151
102	2 Related Work	et al., 2022; Adilazuarda et al., 2024). In reproduc-	152
103	We cluster the closest work into three lines:	tive health, participatory work with underserved	153
104	women’s-health answer evaluation, multilingual	women shows culturally appropriate responses re-	154
105	and sociocultural language understanding, and eval-	quire local context beyond literal translation (Deva	155
106	uation design for ambiguity and safety.	et al., 2025), while Western-centric rubrics can un-	156
107	2.1 Women’s-Health Answer Evaluation	derrate responses trained annotators judge appro-	157
108	Women’s-health LLM research primarily asks	priate (Dey et al., 2025). LUHME work identifies	158
109	whether generated answers are accurate, readable,	related problems: models compress cultural varia-	159
110	empathetic, or safe (Recker et al., 2025; de Vries	tion (Mohammadi et al., 2025), adaptation can fail	160
111	et al., 2025). Adhikary et al. (2025) introduce Men-	when sociocultural context is ignored (Freitas and	161
112	stLLaMA, one of the closest domain-adaptation	Cardoso, 2025), and dialogue systems need mech-	162
113	systems, fine-tuning LLaMA-3-8B on menstrual-	anisms for common ground (Anikina et al., 2025).	163
114	health QA data and evaluating educational response	Concept-level terminological precision, not surface	164
115	quality in English. Mughal et al. (2025) develop	fluency, matters most where domain mismatches	165
116	Mai, a menstrual-health dialogue system in English	can alter meaning, as in healthcare (Heinisch, 2025).	166
117	and Roman Urdu motivated by stigma and socio-	Cross-register transfer work further motivates hold-	167
118	cultural barriers, but evaluated through automatic	ing semantic content constant when measuring sen-	168
119	metrics and user ratings rather than matched inter-	sitivity to linguistic form (Nicholls and Alperin,	169
120	pretation tests. For PCOS, Devranoglu et al. (2024)	2025). Our setting brings these concerns together	170
121	evaluate ChatGPT on clinician-authored questions	in matched English, Arabic, and French women’s-	171
122	and Graça et al. (2026) run an international blinded	health descriptions rather than testing general med-	172
123	survey of health-professional ratings; a systematic	ical knowledge alone.	173
124	review of 80 PCOS AI studies finds the broader field	2.3 Gap and Evaluation Requirements	174
125	still dominated by imaging and structured predictive	The resulting gap is a framework that separates un-	175
126	modeling (Ghaderzadeh et al., 2025), so conversa-	derstanding the input from generating an appealing	176
127	tional LLM evaluation is a small slice of that litera-	answer. Clarification research treats detecting the	177
128	ture. For fertility, Chervenak et al. (2023) compare	need to ask, selecting an informative question, and	178
129	ChatGPT counseling with public-health sources,	responding after clarification as distinct capabilities	179




Figure 1: Conceptual overview. Clinical content is held constant while linguistic form and language vary, separating interpretation failures from changes in the underlying case.

that current models still struggle with (Testoni and Fernández, 2024; Zhang and Choi, 2025); structured, stateful dialogue that elicits missing information turn by turn is one practical direction (Demarco et al., 2024). Such a framework must (i) hold the underlying concern constant across register and language, (ii) score interpretation, risk/action, and clarification need separately, (iii) expose under-triage and unsafe advice rather than rely on aggregate accuracy, and (iv) prevent overlap between evaluation and adaptation data. Native-speaker review targets semantic and register preservation in Arabic and French, while source- and paraphrase-family exclusions separate evaluation from training. The blinded clinical-rating designs used in PCOS and fertility studies motivate independent review of response quality (Graça et al., 2026; Gély et al., 2026), but our primary measures remain tied to predefined item labels.

LLM judges show only moderate agreement and systematic rubric biases (Donati et al., 2025), so rubric-scored cultural sensitivity, helpfulness, clarity, and safety are reported separately from deterministic label metrics; pairing an LLM judge with a human-verification step is one proposed mitigation (Grießhaber et al., 2025).

3 Methods

This is an initial-exploration study: the protocol below describes a general procedure rather than a finished, fixed instrument. Design and construction are described qualitatively; exact item and split counts are intentionally omitted (§5.1).

3.1 Design philosophy: matched cases, controlled register

The evaluation is built around one idea: express a single underlying women’s-health concern in several ways a real patient might, while keeping the clinical facts fixed. Each case is written once in a canonical, evidence-anchored form and rewritten into five further *registers* — clinical, lay, indirect/culturally hedged, ambiguous/under-specified, and emotionally concerned — so that any change




Figure 2: Data construction pipeline (schematic).

in predicted concern category, risk level, or clarification need can be attributed to *how* the case is phrased (Figure 1, Figure 2). Every case carries a predefined target for category, risk, action, and clarification need, so scoring never compares free text directly. For example, a case about heavy menstrual bleeding might appear as a clinical chart-note question in its *clinical* register and as a short, worried, under-specified message in its *ambiguous* register; only the ambiguous variant targets clarification, since it is specifically constructed to omit a detail a clinician would need to ask for (§3.4).

3.2 Seed curation

Candidate patient-authored questions for PCOS/hormonal and fertility content were drawn from the HealthCareMagic-100k and iCliniq-10k patient–doctor dialogue datasets (Li et al., 2023; Li, 2023), accessed via a Hugging Face repackaging (Malikeh1375, 2023), supplemented by additional patient-question corpora for menstrual-health content. Candidates were filtered by keyword for the three categories; cleaned of a recurring text-substitution artifact present in a subset of ChatDoctor-derived answers; filtered for authentic first-person framing against near-duplicate, myth-only, or off-topic content; and balanced across categories. Where a case’s claim mapped to a public-health-authority page, its category was anchored to that source — NHS and NICHD guidance broadly,

Study	Primary target	Languages	Matched register variants	Explicit triage / clarification targets
MenstLLaMA (Adhikary et al., 2025)	Educational answer generation	English	No	No
Mai (Mughal et al., 2025)	Menstrual-health dialogue generation	English, Roman Urdu	No	No
WHBench (Maurya et al., 2026)	Expert-scored clinical responses	English	No	Safety rubric; no clarification target
HerHealthEval (this work)	Symptom-language understanding	English (FR/AR: adaptation data)	Yes	Yes

Table 1: Positioning against the closest women’s-health LLM systems. French and Arabic here refer to the adaptation-data languages; the scored evaluation set is English (§5.1).

with CDC guidance used specifically for fertility content where NHS/NICHD coverage was thin; unmapped claims were flagged rather than assigned an invented category. This anchoring pass is deterministic and rule-based, not model-generated.

3.3 Register construction

Each canonical text is rewritten into the remaining five registers under a meaning-preservation constraint: vocabulary, directness, hedging, and affect may vary, but no clinical claim may be added or removed. An early template-based rewriter collapsed distinct cases onto near-identical clinical-register text and was replaced with a pass that verifies distinctness and meaning preservation. Every register variant is flagged for human spot-check.

3.4 Gold risk and clarification targets

An early approach tried to ground the risk-level target in the same evidence pages used for category anchoring; this was abandoned once it produced a single, non-discriminating label for effectively every case. The risk target used here instead comes from a disclosed, deterministic keyword rule over the case’s answer text, computed once per case so all of its register variants share one ground-truth risk label. The clarification target is register-determined by construction: *yes* exactly for a case’s ambiguous-register row, *no* otherwise. Both targets are explicitly *silver* — heuristic, not clinician-adjudicated — and used identically for the evaluation set and the adaptation data.

3.5 Multilingual localization

Matched English cases are localized into French and Modern Standard Arabic (MSA) rather than a regional dialect, so terminology generalizes across Arabic-speaking audiences. Both corpora were produced with the same LLM translation pipeline (OpenAI’s Responses API, gpt-5.6-sol, low rea-




Figure 3: Model adaptation lineage (schematic).

soning effort, structured outputs), preserving each row’s register label rather than flattening every register into one translation. Output then passes an automated linguistic screen (structural integrity, encoding, source–target alignment) and native-speaker confirmation of meaning and register; for Arabic this is a professional Arabic↔English medical translator working in MSA, checking terminology and grammar (including gender agreement in patient-directed phrasing).

3.6 Evaluation/adaptation separation

Evaluation items are held out from every adaptation corpus by two keys: exact (case-insensitive) question match, and shared source-provenance identity, which also removes paraphrase-family overlap. The rule applies uniformly across English, French, and Arabic, and train/validation duplicates are resolved so the same case cannot inform both.

3.7 Models and adaptation protocol

All experiments use Qwen3.5-9B. We report a zero-shot baseline (M2): the base model prompted with a structured triage schema for category, risk, action, and clarification. We then adapt with QLoRA on the leakage-controlled data above, wrapping targets in the exact structured schema the evaluation items score against (M3); alternative target formats were explored during development and are not reported here. A joint English+French+Arabic run follows the same protocol and is scored the same way (Figure 3); results are reported separately once

available.

3.8 Metrics

We score four predefined-label dimensions rather than free-text quality: *category accuracy*; *risk accuracy*, reported against the majority-class baseline since labels are skewed; *under/over-triage rate*, the safety-relevant direction of risk error; and *clarification recall/specificity*. We also report *parse compliance* and a *cross-style consistency* rate (the fraction of a case’s register variants given the same label); an analogous cross-language rate applies once a scored multilingual item set exists. A model-self-reported `unsafe_response` flag is logged but treated as weak, non-validated signal. Paired differences between the baseline and each recipe use McNemar’s test on matched per-item correctness.

4 Results

We report results on the English evaluation set for two Qwen3.5-9B variants: the zero-shot baseline and our QLoRA adaptation on structured (JSON) targets. Results for French and Arabic will be reported separately once available. Gold labels are skewed, so we foreground majority-class baselines alongside accuracy-style numbers and treat per-class recall, under-triage, and clarification recall as the honest headlines. The risk majority-class baseline is 0.644 and the clarification majority-class baseline is 0.833.

4.1 Zero-shot baseline (M2)

Table 2 reports the zero-shot baseline (M2), an instruction-tuned Qwen3.5-9B prompted with our structured triage schema. M2 parses at 100% and reaches 0.667 risk accuracy, only marginally above the 0.644 majority baseline. The safety-critical weakness is *under-triage*: of 174 items whose gold risk is *see-doctor*, M2 routes 71.8% to routine (per-class recall 0.264, versus 0.897 for routine). M2 retains strong clarification-seeking, however: recall 0.856 on the 90 items requiring clarification, at 0.862 specificity. Category accuracy is 0.539, with fertility weakest (recall 0.467). The model gives a consistent risk label to only 64.4% of case groups across their six phrasings, and a consistent category to just 15.6% — direct evidence that surface communication style, not clinical content, drives much of the interpretation variance (RQ1, §5.1).

Metric	M2	M3	Base
parse_ok	1.000	0.998	—
risk accuracy	0.667	0.623	0.644
under-triage ↓	0.718	0.647	—
clarif. recall	0.856	0.044	—
category accuracy	0.539	0.549	—
consistency (risk)	0.644	0.433	—
consistency (cat.)	0.156	0.344	—

Table 2: Results on the English evaluation set ($n = 540$). M2 = zero-shot baseline; M3 = QLoRA fine-tune on structured (JSON) targets. “Base” = majority-class baseline. Arrows: lower is better for under-triage. Best per row in **bold**. Consistency = fraction of case groups given the same label across all six styles.

4.2 Fine-tuning reduces under-triage but erases clarification-seeking

We fine-tune Qwen3.5-9B with QLoRA on leakage-controlled adaptation data drawn from the same case pool as the evaluation set but disjoint from it, using targets wrapped in the exact structured schema the evaluation items score against.

M3 improves parse compliance (0.998) and reduces under-triage below the zero-shot baseline (0.718 to 0.647), and more than doubles cross-style category consistency (0.156 to 0.344). But its risk accuracy is significantly worse than the zero-shot baseline (0.623 vs. 0.667; McNemar’s $p=0.03$, favoring M2), and clarification recall — 0.856 for the zero-shot baseline on the 90 items that require it — collapses to 0.044: fine-tuning all but eliminates clarification-seeking. Category ($p=0.71$) and clarification ($p=0.30$) correctness are not significantly different from M2; parse compliance is statistically unchanged ($p=1.0$).

5 Discussion

Two results stand out. General-purpose Qwen3.5-9B under-triages the majority of care-warranting cases and interprets the same case differently depending only on phrasing (RQ1, §5.1). Adaptation is a double-edged tool: it reduces under-triage and stabilizes category interpretation, but at some cost to risk accuracy, and it erases clarification-seeking almost entirely. This suggests clarification on ambiguous input is not a skill that answer-style supervised fine-tuning teaches by example; recovering it likely needs objectives that explicitly reward abstention or asking — preference optimization, or a dedicated clarification-vs-answer decision head — rather than more demonstrations. The practical lesson for safety-relevant domains: evaluate adaptation on the structured behaviors that matter at

deployment, triage *and* clarification, not answer quality alone, since a model can improve on one while losing another.

5.1 Preliminary answers to the research questions

We revisit the three research questions against the evidence available at this stage; all are answered only as far as current results allow.

RQ1 (register robustness). *Partially answered, English only.* Holding a case fixed and varying only its phrasing changes the predicted risk label for 35.6% of cases and the predicted category for 84.4% (cross-style consistency 0.644 and 0.156). Adaptation partly redistributes this: M3 more than doubles category consistency (0.156 to 0.344) but lowers risk consistency (0.644 to 0.433), so robustness gains are not uniform across target types.

RQ2 (cross-lingual differences). *Not yet answerable.* Results for Arabic and French will be reported separately once available; we do not extrapolate the English-only findings above to other languages.

RQ3 (adaptation impact). *Partially answered, English-only adaptation.* Domain adaptation reduces under-triage (0.718 to 0.647) and improves cross-style consistency, but at a cost to risk accuracy and with clarification recall collapsing to near zero. Results for the multilingual (English+French+Arabic) adaptation stage will be reported separately once available.

Limitations

Results reported here are on the English evaluation set; results for French and Arabic will follow separately. Gold labels are skewed toward routine risk and no-clarification, which is why we foreground majority baselines, per-class recall, and under-triage over raw accuracy. Risk and clarification targets are a disclosed heuristic rather than clinician adjudication, applied identically to the evaluation set and the adaptation data (§3.4). The `unsafe_response` flag is model-self-reported, not an independently validated safety measure. The Arabic adaptation data uses Modern Standard Arabic only. Since Qwen3.5-9B is itself a multilingual model with pretrained exposure to Arabic dialects, this is a gap in our controlled data coverage rather than evidence that the model cannot handle dialectal input — but we have not directly measured dialectal performance, and extending evaluation to dialectal

Arabic (e.g., Egyptian, Gulf, Levantine) is a natural direction for future work. This initial study covers three symptom categories, three languages, and a single 9B model family, so generalization to other conditions, cultures, care settings, and model scales is untested. Arabic and French adaptation text was produced by LLM machine translation and confirmed by native-speaker/professional review for meaning and register (§3); this validates translation quality, not the clinical accuracy of gold targets. Model outputs are evidence of linguistic understanding in a controlled research setting, not medical advice or proof of clinical-deployment readiness.

References

- Prattay Kumar Adhikary, Isha Motiyani, Gayatri Oke, Maithili Joshi, Kanupriya Pathak, Salam Michael Singh, and Tanmoy Chakraborty. 2025. [Menstrual health education using a specialized large language model in India: Development and evaluation study of MenstLLaMA](#). *Journal of Medical Internet Research*, 27:e71977.
- Muhammad Farid Adilazuarda, Sagnik Mukherjee, Pradhyumna Lavania, Siddhant Shivdutt Singh, Alham Fikri Aji, Jacki O’Neill, Ashutosh Modi, and Monojit Choudhury. 2024. [Towards measuring and modeling “culture” in LLMs: A survey](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 15763–15784, Miami, Florida, USA. Association for Computational Linguistics.
- Tatiana Anikina, Alina Leippert, and Simon Ostermann. 2025. [Building common ground in dialogue: A survey](#). In *Proceedings of the 2nd LUHME Workshop*, pages 3–28, Bologna, Italy. UP - Universidade do Porto (<https://doi.org/10.21747/978-989-9193-73-4/lan2>), LIACC - Laboratório de Inteligência Artificial e Ciência de Computadores da Universidade do Porto, CLUP - Centro de Linguística da Universidade do Porto, UEF - The University of Eastern Finland and UAH - Universidad de Alcalá.
- Joseph Chervenak, Harry Lieman, Miranda Blanco-Breindel, and Sangita Jindal. 2023. [The promise and peril of using a large language model to obtain clinical information: ChatGPT performs strongly as a fertility counseling tool with limitations](#). *Fertility and Sterility*, 120(3):575–583.
- P. L. M. de Vries, D. Baud, S. Baggio, M. Ceulemans, G. Favre, E. Gerbier, et al. 2025. Enhancing perinatal health patient information through ChatGPT—an accuracy study. *PEC Innovation*, 6:100381.
- Gianfranco Demarco, Nicola Fanelli, Gennaro Vessio, and Giovanna Castellano. 2024. [Converso: Improving LLM chatbot interfaces and task execution via conversational form](#). In *Proceedings of the 1st*

504	LUHME Workshop, pages 5–11, Santiago de Compostela, Spain. CLUP, Centro de Linguística da Universidade do Porto FLUP - Faculdade de Letras da Universidade do Porto.	559
505		560
506		561
507		562
		563
508	Roshini Deva, Dhruv Ramani, Tanvi Divate, Suhani Jalota, and Azra Ismail. 2025. “Kya Family Planning after Marriage Hoti Hai?”: Integrating cultural sensitivity in an LLM chatbot for reproductive health. In <i>Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems</i> .	564
509		565
510		566
511		567
512		568
513		569
514	Belgin Devranoglu, Tugba Gurbuz, and Oya Gokmen. 2024. ChatGPT’s efficacy in queries regarding polycystic ovary syndrome and treatment strategies for women experiencing infertility. <i>Diagnostics</i> , 14(11):1082.	570
515		571
516		572
517		573
518		574
519	Sumon Kanti Dey, Manvi S, Zeel Mehta, Meet Shah, Unnati Agrawal, Suhani Jalota, and Azra Ismail. 2025. Beyond the rubric: Cultural misalignment in LLM benchmarks for sexual and reproductive health. In <i>Proceedings of the 5th Workshop on Evaluation and Comparison of NLP Systems</i> , pages 126–134, Mumbai, India. Association for Computational Linguistics.	575
520		576
521		577
522		578
523		579
524		580
525		581
		582
526	Nicolò Donati, Paolo Torroni, and Giuseppe Savino. 2025. Do large language models understand how to be judges? In <i>Proceedings of the 2nd LUHME Workshop</i> , pages 85–102, Bologna, Italy. UP - Universidade do Porto (https://doi.org/10.21747/978-989-9193-73-4/lan2), LIACC - Laboratório de Inteligência Artificial e Ciência de Computadores da Universidade do Porto, CLUP - Centro de Linguística da Universidade do Porto, UEF - The University of Eastern Finland and UAH - Universidad de Alcalá.	583
527		584
528		
529		585
530		586
531		587
532		588
533		589
534		590
535		591
		592
536	David C. T. Freitas and Henrique Lopes Cardoso. 2025. A nightmare on LLMs street: On the importance of cultural awareness in text adaptation for LRLs. In <i>Proceedings of the 2nd LUHME Workshop</i> , pages 40–48, Bologna, Italy. UP - Universidade do Porto (https://doi.org/10.21747/978-989-9193-73-4/lan2), LIACC - Laboratório de Inteligência Artificial e Ciência de Computadores da Universidade do Porto, CLUP - Centro de Linguística da Universidade do Porto, UEF - The University of Eastern Finland and UAH - Universidad de Alcalá.	593
537		594
538		
539		595
540		596
541		597
542		598
543		599
544		
545		600
546		601
547	Lilas Gély, Maxime Chaillot, and Thomas Fréour. 2026. Can large language models provide accurate and empathetic answers to the most frequently asked questions by infertile patients? a pilot study. <i>Reproductive BioMedicine Online</i> , 52(4):105221.	602
548		603
549		604
550		605
551		606
552	Mustafa Ghaderzadeh, Ali Garavand, and Cirruse Salehnasab. 2025. Artificial intelligence in polycystic ovary syndrome: A systematic review of diagnostic and predictive applications. <i>BMC Medical Informatics and Decision Making</i> , 25:427.	607
553		608
554		
555		609
556		610
		611
		612
557	Sandro Graça, Alexander Dallaway, Folashade Alloh, Harpal S. Randeva, Chris Kite, and Ioannis Kyrou. 2026. Assessing ChatGPT vs. evidence-based online responses for polycystic ovary syndrome self-management and education: An international cross-sectional blinded survey of healthcare professionals. <i>Frontiers in Digital Health</i> , 7:1700018.	613
558		614
		615
		616
		617
		618
		619
		620
		621
		622
		623
		624
		625
		626
		627
		628
		629
		630
		631
		632
		633
		634
		635
		636
		637
		638
		639
		640
		641
		642
		643
		644
		645
		646
		647
		648
		649
		650
		651
		652
		653
		654
		655
		656
		657
		658
		659
		660
		661
		662
		663
		664
		665
		666
		667
		668
		669
		670
		671
		672
		673
		674
		675
		676
		677
		678
		679
		680
		681
		682
		683
		684
		685
		686
		687
		688
		689
		690
		691
		692
		693
		694
		695
		696
		697
		698
		699
		700
		701
		702
		703
		704
		705
		706
		707
		708
		709
		710
		711
		712
		713
		714
		715
		716
		717
		718
		719
		720
		721
		722
		723
		724
		725
		726
		727
		728
		729
		730
		731
		732
		733
		734
		735
		736
		737
		738
		739
		740
		741
		742
		743
		744
		745
		746
		747
		748
		749
		750
		751
		752
		753
		754
		755
		756
		757
		758
		759
		760
		761
		762
		763
		764
		765
		766
		767
		768
		769
		770
		771
		772
		773
		774
		775
		776
		777
		778
		779
		780
		781
		782
		783
		784
		785
		786
		787
		788
		789
		790
		791
		792
		793
		794
		795
		796
		797
		798
		799
		800
		801
		802
		803
		804
		805
		806
		807
		808
		809
		810
		811
		812
		813
		814
		815
		816
		817
		818
		819
		820
		821
		822
		823
		824
		825
		826
		827
		828
		829
		830
		831
		832
		833
		834
		835
		836
		837
		838
		839
		840
		841
		842
		843
		844
		845
		846
		847
		848
		849
		850
		851
		852
		853
		854
		855
		856
		857
		858
		859
		860
		861
		862
		863
		864
		865
		866
		867
		868
		869
		870
		871
		872
		873
		874
		875
		876
		877
		878
		879
		880
		881
		882
		883
		884
		885
		886
		887
		888
		889
		890
		891
		892
		893
		894
		895
		896
		897
		898
		899
		900
		901
		902
		903
		904
		905
		906
		907
		908
		909
		910
		911
		912
		913
		914
		915
		916
		917
		918
		919
		920
		921
		922
		923
		924
		925
		926
		927
		928
		929
		930
		931
		932
		933
		934
		935
		936
		937
		938
		939
		940
		941
		942
		943
		944
		945
		946
		947
		948
		949
		950
		951
		952
		953
		954
		955
		956
		957
		958
		959
		960
		961
		962
		963
		964
		965
		966
		967
		968
		969
		970
		971
		972
		973
		974
		975
		976
		977
		978
		979
		980
		981

language models understand morality across cultures? In *Proceedings of the 2nd LUHME Workshop*, pages 30–39, Bologna, Italy. UP - Universidade do Porto (<https://doi.org/10.21747/978-989-9193-73-4/lan2>), LIACC - Laboratório de Inteligência Artificial e Ciência de Computadores da Universidade do Porto, CLUP - Centro de Linguística da Universidade do Porto, UEF - The University of Eastern Finland and UAH - Universidad de Alcalá.

Shafaq Fatima Mughal, Shayan Aamir, Abdul Samad, Umema Zehra, and Ayesha Asif Syed. 2025. *Mai: A transformer based domain specific chatbot for menstrual health*. *ACM Journal on Responsible Computing*, 2(1):1–32.

Robin Nicholls and Kenneth Alperin. 2025. *Cross-genre native language identification with open-source large language models*. In *Proceedings of the 2nd LUHME Workshop*, pages 103–108, Bologna, Italy. UP - Universidade do Porto (<https://doi.org/10.21747/978-989-9193-73-4/lan2>), LIACC - Laboratório de Inteligência Artificial e Ciência de Computadores da Universidade do Porto, CLUP - Centro de Linguística da Universidade do Porto, UEF - The University of Eastern Finland and UAH - Universidad de Alcalá.

Florian Recker, Ricarda Neubauer, Agnes Wittek, and Nadine Scholten. 2025. *Large language models and women’s health: a digital companion for informed decision-making*. *Archives of Gynecology and Obstetrics*, 312:663–670.

Ipek Baris Schlicht, Zhixue Zhao, Burcu Sayin, Lucie Flek, and Paolo Rosso. 2025. *Do LLMs provide consistent answers to health-related questions across languages?* *Preprint*, arXiv:2501.14719.

Alberto Testoni and Raquel Fernández. 2024. *Asking the right question at the right time: Human and model uncertainty guidance to ask clarification questions*. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 258–275, St. Julian’s, Malta. Association for Computational Linguistics.

Michael JQ Zhang and Eunsol Choi. 2025. *Clarify when necessary: Resolving ambiguity through interaction with LMs*. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 5541–5558, Albuquerque, New Mexico. Association for Computational Linguistics.